

Advanced Graph Theory Topics

Andre Kessler

December 3, 2010

1 Tree Tricks

Given a connected graph G with V vertices and $V - 1$ edges, find the maximum distance between any two nodes. How can we do this in linear time? We make the following important observation: if X is the node in G that is the farthest from the root, the distance from X to the farthest node in the graph is the maximum distance between *any* two graph nodes. To prove this, we will argue by contradiction. Suppose there exist two nodes $A, B \in G$ with the distance between A and B greater than any between X and some other node. Then let L be the least common ancestor of A and B . Since X is the deepest node in the party, the least common ancestor of A and X must be below L , because A is closer to X than B . The same argument holds for B , but this is a contradiction, as it implies there is a cycle within the tree. It is important to be able to recognize when a problem uses a tree as opposed to a generic graph, as many graph algorithms will vastly simplify. Thus, you should be able to do this even when the conditions are weird. A list of some necessary and sufficient conditions is below.

- G is connected and acyclic.
- G is connected and has $V - 1$ edges.
- G has no cycles and has $V - 1$ edges.
- If any edge is added to G , we get a cycle.
- If any edge is removed from G , the graph becomes disconnected.
- Any two vertices in G are connected by a unique path.

2 Minimal Spanning Trees

Given a connected graph G with weighted edges, remove edges so that the graph has minimum total weight possible but remains connected. How might we go about finding this tree? More formally, suppose E is the set of edges in G . Then we want to find an acyclic subset $T \subset E$ that connects all of the vertices and whose total weight $w(T)$ is minimized. This set of edges is acyclic and contains all the vertices, so it must be a tree. In this case we call it a *spanning tree* since its span covers G . For an example of this, take a look¹ at Figures 1 and 2.

2.1 Generic Algorithm

First, we will describe a generic algorithm that uses a greedy strategy. The algorithm keeps a current set of edges T while keeping the following invariant: prior to each iteration, T is a subset of some minimum spanning tree. In each step of our algorithm, we will need to determine an edge e that can be added to T which is such that $T \cup \{e\}$ is also a subset of some minimum spanning tree. Let's cut the graph (partition it into two sets). Suppose our cut *respects* the set T , meaning it does not cut through any edges in T . Then take an edge $e_1 = (u, v)$ of minimum weight crossing the cut. Adding this to T clearly does not form a cycle. Additionally, any minimum spanning tree

¹Diagram from the USACO training pages.




Figure 1: Our original graph G .




Figure 2: One minimum spanning tree.

of the graph G must have an edge crossing the cut. Suppose this is some different edge; call it $e_2 = (x, y)$. At the current point in time, e_2 is not in T , because our cut respects T . Since e_2 is on the unique path from u to v in the minimal spanning tree containing our current set T , removing e_2 splits T into two components. Adding edge $e_1 = (u, v)$ reconnects the graph, and does not increase the weight. Hence this edge must be safe, and our algorithm works.

2.2 Prim's Algorithm

Once armed with our generic algorithm, Prim's Algorithm is not difficult to visualize and understand. We simply start with an arbitrary root vertex and grow this single tree until we can't add any more edges. At each point in the algorithm, our "cut" will be the frontier of the current set T . Adding the smallest edge here preserves the required invariant and so results in a minimal spanning tree. Stated more simply, we start with a tree containing any one node of our graph G . Iteratively, we find the closest node to that one and add the edge between them. Now we continue to build our tree. At each step, we find a node not currently within the tree which is closest to the tree and add the minimum weight edge from that node to some node in the tree, and add it to the tree.

My example C++ code is below:

```

1  int prim (int s)
2  {
3      int weight = 0; bool vis [MAXV];
4      memset (vis, false, sizeof (vis));
5      priority_queue <state> q;
6      q.push ((state) {0, s});
7
8      while (!q.empty ())
9      {
10         state top = q.top (); q.pop ();
11
12         if (!vis [top.node])
13         {
14             vis [top.node] = true;
15             weight += top.dist;
16
17             for (int i = 0; i < adj [top.node].size (); i++)
18                 q.push ((state) {adj [top.node][i].dist, adj [top.node][i].node});
19         }
20     }
21
22     return weight;
23 }

```

Notice how I use a priority queue to keep track of the “frontier” nodes and process them in order of weight. Careful readers will realize that this is virtually identical to the implementation of Dijkstra’s algorithm. This also immediately implies that our complexity is $O(E \log V)$.

2.3 Kruskal’s Algorithm

I won’t go into much detail on Kruskal’s algorithm right now because it involves disjoint-set data structures, which will be the topic of a completely different lecture. However, the idea behind it is important. Kruskal’s algorithm is almost directly based on the generic MST algorithm described earlier. We iterate through a sorted list of edges, taking all edges in that order which are “safe” (that is, do not form a cycle). This checking can be done very quickly with a disjoint-set structure, leading to an overall runtime of $O(E\alpha(V))$, where $\alpha(V)$ is the incredibly slow-growing inverse Ackermann function. This is so slow that $\alpha(V) \leq 5$ for all practical V . In fact, the bottleneck becomes the sort which is necessary if the edges are not given in sorted order; in this case, the complexity is $O(E \log V)$.

3 Shortest Paths: the Next Step

You’ve done breadth-first-search, Dijkstra, and Floyd-Warshall. What’s next? Now we take a look at shortest path algorithms for more specific situations. There are many cases in which shortest path algorithms simplify for certain graphs. For example, in a tree, we can compute the distance between all pairs of nodes in $O(N^2)$ time using the formula

$$\text{dist}(A, B) = \text{depth}(A) + \text{depth}(B) - 2 \cdot \text{depth}(\text{LCA}(A, B)).$$

In directed acyclic graphs, or DAGs, linear time algorithms for shortest path are pretty easy. Additionally, problems sometimes crop up involving negative edge weights on a graph or detection of

negative cycles, with the latter often occurring as a subproblem. For this, we want the Bellman-Ford algorithm. But before we dive in, let's quickly review the theoretical underpinning of shortest-path algorithms: relaxation.

3.1 Relaxation

For each vertex $v \in G$, we will keep a value $d(v)$, which is an upper bound on the weight of the shortest path from the source to node v . Initially, $d(\text{source}) = 0$ and $d(v)$ for all other v is set to infinity. When we relax edge (u, v) , we replace $d(v)$ by $d(u) + w(u, v)$ (if necessary).

3.2 Shortest Paths in a DAG

Directed acyclic graphs are nice because shortest paths all from a source are very easy to compute in linear time. First, if this is not already the case, we topologically sort the DAG. Then, we make just one pass over the vertices in the topological order, relaxing the edges that leave the vertex we're on as we go.

3.3 Bellman-Ford

We use the Bellman-Ford algorithm when we want to find shortest paths in a graph where some edges may have negative weight or determine if a graph contains a negative cycle. This algorithm is also somewhat unusual in that we take an edge list for input. Note that a shortest path will not contain more than $V - 1$ edges in the graph, assuming there is no negative cycle. If there are more than $V - 1$ edges in the shortest path, some node will have been visited twice, which is not optimal. My sample code is below.

```
1 // Initialization
2 for (int i = 0; i < V; i++)
3     dist [i] = INFTY;
4
5 dist [0] = 0;
6
7 // Bellman-Ford Algorithm
8 for (int i = 1; i < V; i++)
9     for (int j = 0; j < edges.size (); j++)
10         if (dist [edges [j].b] > dist [edges [j].a] + edges [j].len)
11             dist [edges [j].b] = dist [edges [j].a] + edges [j].len;
12
13 // Negative-cycle detection
14 for (int i = 0; i < E; i++)
15     if (dist [edges [i].b] > dist [edges [i].a] + edges [i].len)
16     {
17         puts ("Oh no, a negative cycle!");
18         break;
19     }
```

This clearly runs in time $O(EV)$, so it is generally worse than Dijkstra's, although it is easier to code. However, you generally should only use Bellman-Ford when you specifically need to deal with negative-weight edges.

4 Problems

1. Suppose the graph for Prim's algorithm is represented as an adjacency matrix. Give a simple implementation of Prim's algorithm that runs in $O(V^2)$ time.
2. A minimum bottleneck spanning tree is a spanning tree which minimizes the maximum edge weight over all such trees. Find one such minimum bottleneck spanning tree in a graph G .
3. Give an algorithm for finding the spanning tree with the smallest product of edge weights, assuming all edge weights are positive.
4. Given a graph G and spanning tree T . Suppose we decrease the weight of one of the edges not in T . Find the minimum spanning tree in the modified graph.
5. Bessie the cow is walking through some very corrupt cities in Zimbabwe. The cities are connected by a set of roads, and each road has some bandit associated with it. Some bandits are stronger than her, and will require a specific fee $-X_i$ for one safe passage. Other bandits are weaker than Bessie, so she can force them to give her up to X_i amount of money each time she traverses the road. Please help her compute the maximum amount of money she can obtain on her journey from Harare to Kwekwe.
6. (IOI 2003) Farmer John's cows wish to travel freely among the N ($1 \leq N \leq 200$) fields (numbered $1..N$) on the farm, even though the fields are separated by forest. The cows wish to maintain trails between pairs of fields so that they can travel from any field to any other field using the maintained trails. Cows may travel along a maintained trail in either direction. The cows do not build trails. Instead, they maintain wild animal trails that they have discovered. On any week, they can choose to maintain any or all of the wild animal trails they know about. Always curious, the cows discover one new wild animal trail at the beginning of each week. They must then decide the set of trails to maintain for that week so that they can travel from any field to any other field. Cows can only use trails which they are currently maintaining. The cows always want to minimize the total length of trail they must maintain. The cows can choose to maintain any subset of the wild animal trails they know about, regardless of which trails were maintained the previous week. Wild animal trails (even when maintained) are never straight. Two trails that connect the same two fields might have different lengths. While two trails might cross, cows are so focused, they refuse to switch trails except when they are in a field. At the beginning of each week, the cows will describe the wild animal trail they discovered. Your program must then output the minimum total length of trail the cows must maintain that week so that they can travel from any field to any other field, if there exists such a set of trails.
7. (APIO 2008) The Kingdom of New Asia contains N villages connected by M roads. Some roads are made of cobblestones, and others are made of concrete. Keeping roads free-of-charge needs a lot of money, and it seems impossible for the Kingdom to maintain every road. A new road maintaining plan is needed. The King has decided that the Kingdom will keep as few free roads as possible, but every two distinct villages must be connected by one and only one path of free roads. Also, although concrete roads are more suitable for modern traffic, the King thinks walking on cobblestones is interesting. As a result, he decides that exactly K cobblestone roads will be kept free. Given a description of roads in New Asia and the number of cobblestone roads that the King wants to keep free, write a program to determine if there is a road maintaining plan that satisfies the King's criteria, and output a valid plan if there is one.

8. (David Benjamin and Jacob Steinhardt, 2008) Farmer John has grown so lazy that he no longer wants to continue maintaining the cow paths that currently provide a way to visit each of his N ($5 \leq N \leq 10,000$) pastures (conveniently numbered $1..N$). Each and every pasture is home to one cow. FJ plans to remove as many of the P ($N - 1 \leq P \leq 100,000$) paths as possible while keeping the pastures connected. You must determine which $N - 1$ paths to keep. Bidirectional path j connects pastures S_j and E_j ($1 \leq S_j \leq N; 1 \leq E_j \leq N; S_j \neq E_j$) and requires L_j ($0 \leq L_j \leq 1,000$) time to traverse. No pair of pastures is directly connected by more than one path. The cows are sad that their transportation system is being reduced. You must visit each cow at least once every day to cheer her up. Every time you visit pasture i (even if you're just traveling through), you must talk to the cow for time C_i ($1 \leq C_i \leq 1,000$). You will spend each night in the same pasture (which you will choose) until the cows have recovered from their sadness. You will end up talking to the cow in the sleeping pasture at least in the morning when you wake up and in the evening after you have returned to sleep. Assuming that Farmer John follows your suggestions of which paths to keep and you pick the optimal pasture to sleep in, determine the minimal amount of time it will take you to visit each cow at least once in a day.
9. (Andre Kessler, 2009) Farmer John's cows are living on N ($2 \leq N \leq 200,000$) different pastures conveniently numbered $1..N$. Exactly $N - 1$ bidirectional cow paths (of unit length) connect these pastures in various ways, and each pasture is reachable from any other cow pasture by traversing one or more of these paths. The cows have organized K ($1 \leq K \leq N/2$) different political parties conveniently numbered $1..K$. Each cow identifies with a single political party; cow i identifies with political party A_i ($1 \leq A_i \leq K$). Each political party includes at least two cows. The political parties are feuding and would like to know how much 'range' each party covers. The range of a party is the largest distance between any two cows within that party (over cow paths). Please help the cows determine party ranges.
10. (IOI 2002) Yong-In city plans to build a bus network with N bus stops. Each bus stop is at a street corner. Yong-In is a modern city, so its map is a grid of square blocks of equal size. Two of these bus stops are to be selected as hubs H_1 and H_2 . The hubs will be connected to each other by a direct bus line and each of the remaining $N - 2$ bus stops will be connected directly to either H_1 or H_2 (but not to both), but not to any other bus stop. The distance between any two bus stops is the length of the shortest possible route following the streets. That is, if a bus stop is represented as (x, y) with x -coordinate x and y -coordinate y , then the distance between two bus stops (x_1, y_1) and (x_2, y_2) is $|x_1 - x_2| + |y_1 - y_2|$. If bus stops A and B are connected to the same hub H_1 , then the length of the route from A to B is the sum of the distances from A to H_1 and from H_1 to B . If bus stops A and B are connected to different hubs, e.g., A to H_1 and B to H_2 , then the length of the route from A to B is the sum of the distances from A to H_1 , from H_1 to H_2 , and from H_2 to B . The planning authority of Yong-In city would like to make sure that every citizen can reach every point within the city as quickly as possible. Therefore, city planners want to choose two bus stops to be hubs in such a way that in the resulting bus network the length of the longest route between any two bus stops is as short as possible. One choice P of two hubs and assignments of bus stops to those hubs is better than another choice Q if the length of the longest bus route is shorter in P than in Q . Your task is to write a program to compute the length of this longest route for the best choice P .